

Machine Learning 1


Delaney (PID: A15567985)

2/8/2022

First up kmeans()


Demo of using kmeans() function in base R. First make up some data with a known structure

```
tmp <- c( rnorm(30, -3), rnorm(30, 3) )  
x <- cbind(tmp, y=rev(tmp))  
plot(x)
```



Now we have some made up data in 'x' lets see how kmeans works with this data

```
k <- kmeans(x, centers=2, nstart=20)  
k
```

Now for Hierarchical Clustering

We will cluster the same data 'x' with the `hclust()`. In this case '`hclust()`' requires a distance matrix as input.


```
hc <- hclust(dist(x))
hc
```

```
##
## Call:
## hclust(d = dist(x))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

Let's plot our `hclust` result

```
plot(hc)
```

Cluster Dendrogram



```
dist(x)
hclust (*, "complete")
```


To get our cluster membership vector we need to “cut” the tree with the ‘cutree()’.

```
grps <- cutree(hc, h=8)
grps
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Now plot our data with the `hclust()` results.

```
plot(x, col=grps)
```



Principal Component Analysis (PCA)

PCA of UK food data

Read data from website and try a few visualizations.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
x
```

##	England	Wales	Scotland	N.Ireland
## Cheese	105	103	103	66
## Carcass_meat	245	227	242	267
## Other_meat	685	803	750	586
## Fish	147	160	122	93
## Fats_and_oils	193	235	184	209
## Sugars	156	175	147	139
## Fresh_potatoes	720	874	566	1033
## Fresh_Veg	253	265	171	143
## Other_Veg	488	570	418	355
## Processed_potatoes	198	203	220	187
## Processed_Veg	360	365	337	334
## Fresh_fruit	1102	1137	957	674
## Cereals	1472	1582	1462	1494

## Beverages	57	73	53	47
## Soft_drinks	1374	1256	1572	1506
## Alcoholic_drinks	375	475	458	135
## Confectionery	54	64	62	41

Data Import

Q1. How many rows and columns are in your new data frame named 'x'? What R functions could you use to answer this question?

```
dim(x)
```

```
## [1] 17  4
```

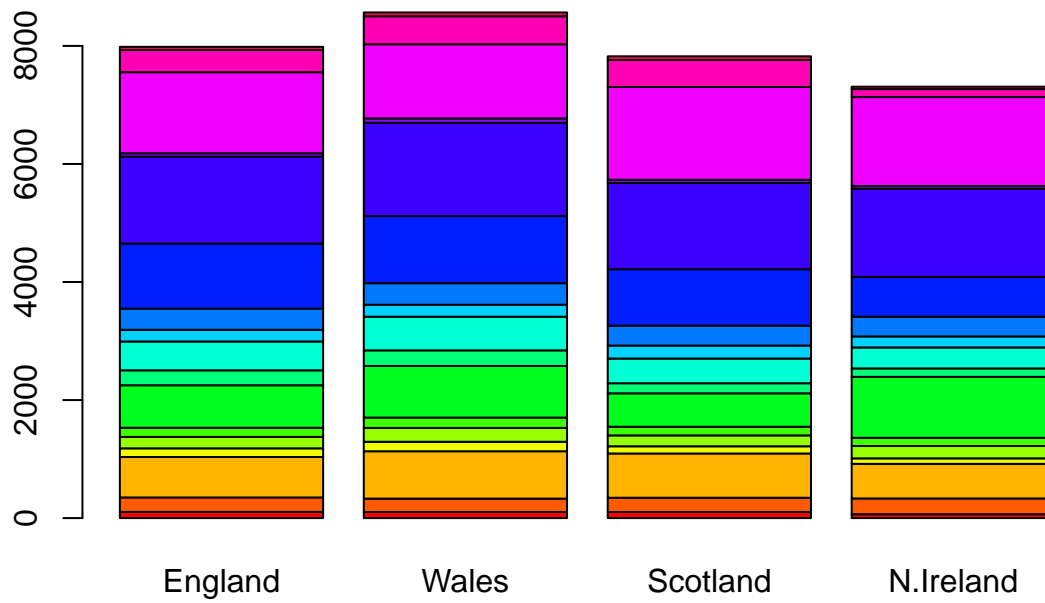
There are 17 rows and 4 columns for the new data frame named 'x'. The R functions that can be used to answer this are 'dim(x)' for both or 'ncol()' and 'nrow()' separately.

Checking your data

Q2. Which approach to solving the 'rownames problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?


I prefer using the argument setting 'row.names()' to set the correct row-names because it only returned the data frame with the row names changed. 'Rownames()' is a non-generic function, while 'row.names()' is a generic function and is specific for data frames.

```
cols <- rainbow(nrow(x))
barplot( as.matrix(x), col=cols )
```



Spotting major differences and trends

```
barplot( as.matrix(x), col=cols, beside=TRUE )
```




Q3. Changing what optional argument in the above `barplot()` function results in the following plot?

Changing `'beside=TRUE'` changes each column to be next to each other rather than stacked.

Q5. Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

The following creates a matrix of scatter plots to help understand the pairwise relationship between the different variables in the data set. If a given point lies on the diagonal for a given plot, it means that the same amount of people eat the same food in the different countries.

```
pairs(x, col= rainbow(10), pch= 16)
```

>Q6. What is the main difference between N. Ireland and the other countries of the UK in terms of this data set?

N. Ireland generally consumes less fresh fruit and eats more potatoes and soft drinks than the other countries of the UK.

PCA to the rescue

The main base R PCA function is called 'prcomp()' and we will need to give it the transpose of our input data!

```
pca <- prcomp( t(x))
summary(pca)
```

```
## Importance of components:
```

```
##              PC1      PC2      PC3      PC4
## Standard deviation 324.1502 212.7478 73.87622 4.189e-14
## Proportion of Variance 0.6744 0.2905 0.03503 0.000e+00
## Cumulative Proportion 0.6744 0.9650 1.00000 1.000e+00
```

```
attributes(pca)
```


```
## $names
## [1] "sdev"      "rotation" "center"    "scale"     "x"
##
```

```
## $class
## [1] "prcomp"
```

To make our new PCA plot (a.k.a. PCA score plot) we access 'pca\$x'.

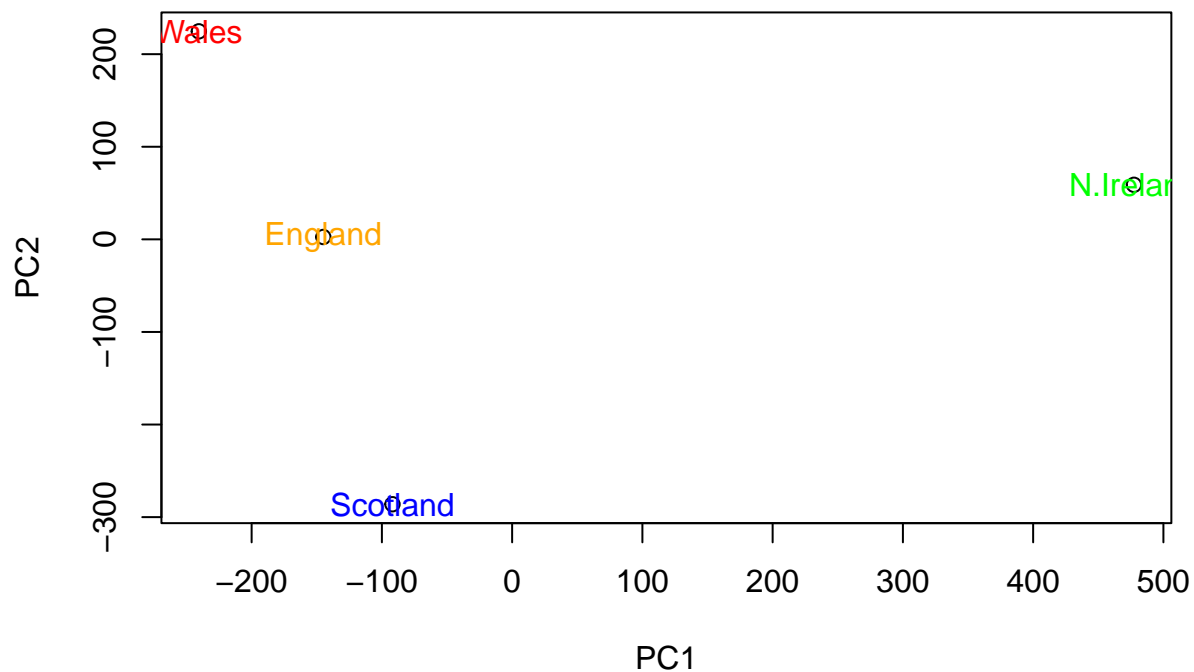
Q7. Complete the code below to get the plot of PC1 vs PC2. The second line adds text labels over the data points.

```
plot(pca$x[,1], pca$x[,2])
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at the start of this document.

```
country_cols <-c("orange", "red", "blue", "green")
plot(pca$x[,1], pca$x[,2], xlab= "PC1", ylab= "PC2")
text(pca$x[,1], pca$x[,2], colnames(x), col= country_cols)
```



Use square of `pca$sdev` to calculate how much variation in the original data each PC accounts to

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

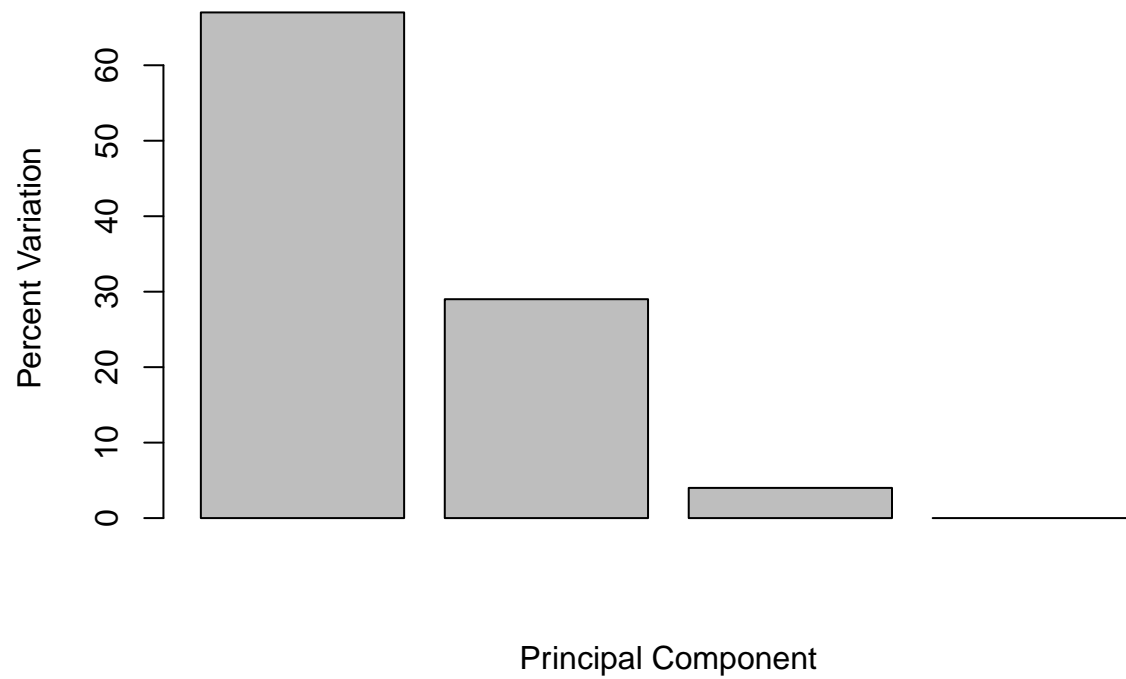
```
## [1] 67 29 4 0
```

Summarize in a plot of variances

```
z <- summary(pca)
z$importance
```

```
##
## Standard deviation      PC1      PC2      PC3      PC4
## Proportion of Variance  0.67444  0.29052  0.03503  0.000000e+00
## Cumulative Proportion   0.67444  0.96497  1.00000  1.000000e+00
```


```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



Digging deeper (variable loadings)

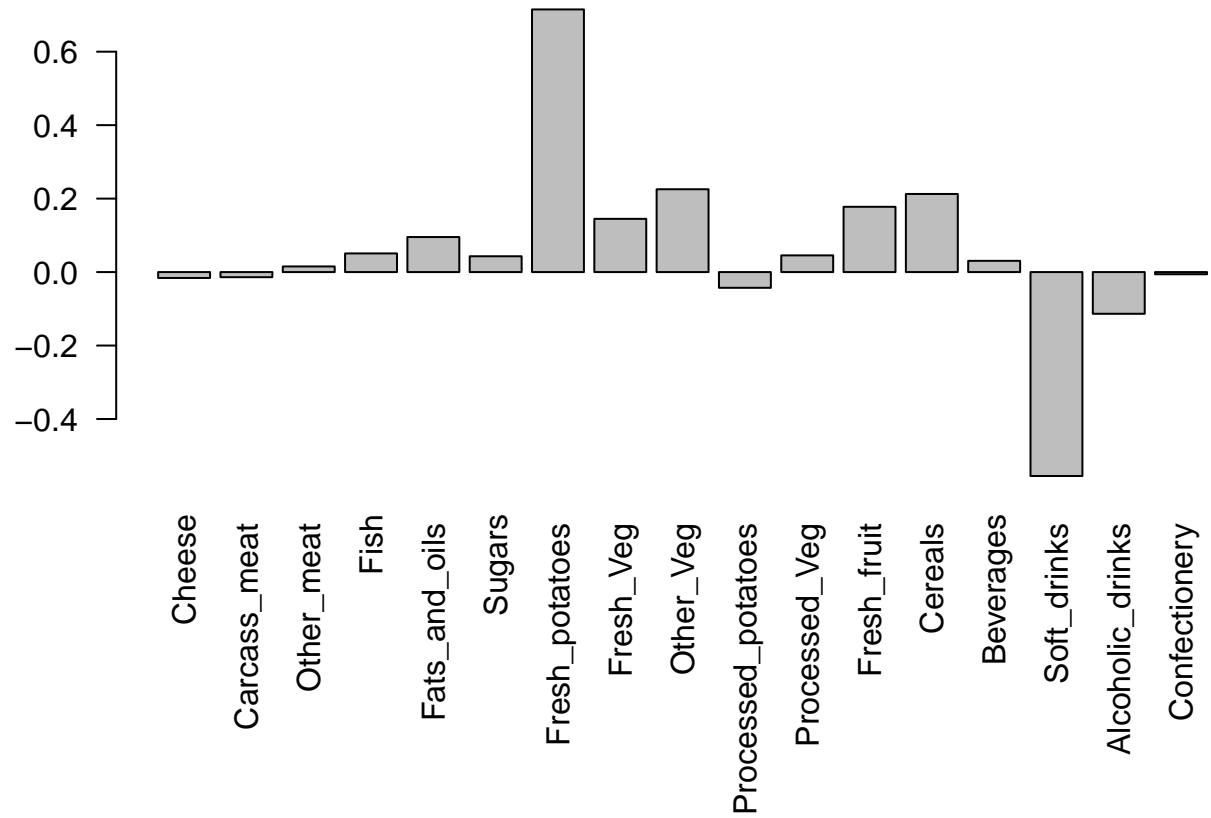
Focus on PC1 as it accounts for >90% of variance

```
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



Q9. Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

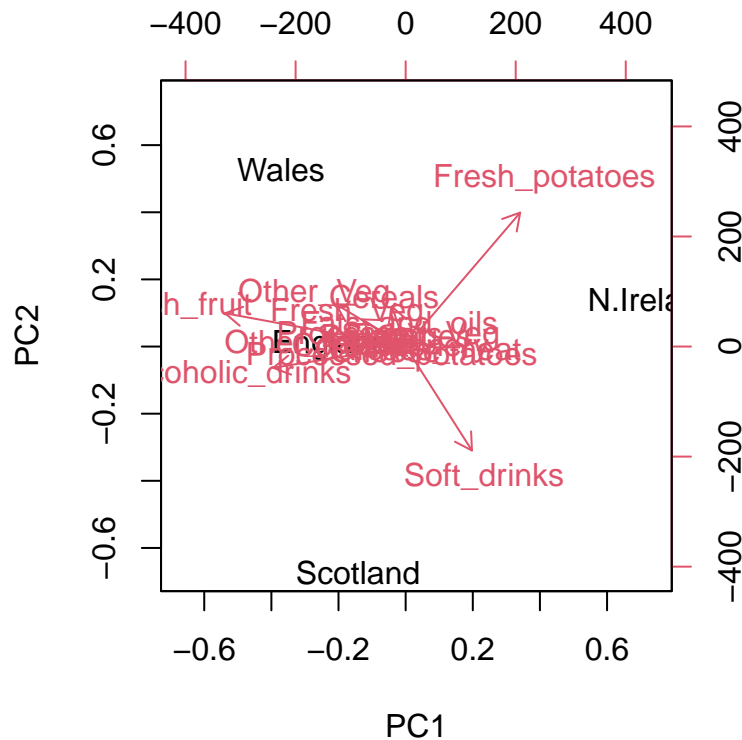
```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```



PC2 mainly shows how N. Ireland eats less fresh fruits and more potatoes and soft drinks.

Biplots

```
biplot(pca)
```



PCR of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

##		wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
##	gene1	439	458	408	429	420	90	88	86	90	93
##	gene2	219	200	204	210	187	427	423	434	433	426
##	gene3	1006	989	1030	1017	973	252	237	238	226	210
##	gene4	783	792	829	856	760	849	856	835	885	894
##	gene5	181	249	204	244	225	277	305	272	270	279
##	gene6	460	502	491	491	493	612	594	577	618	638

Q10. How many genes and samples are in this data set?

```
dim(rna.data)
```

```
## [1] 100 10
```

There are 100 genes and 10 samples in this data set.

Take the transpose of our data

```
pca <- prcomp(t(rna.data), scale=TRUE)
```


This is a summary of how well PCA is doing

```
summary(pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  9.6237 1.5198 1.05787 1.05203 0.88062 0.82545 0.80111
## Proportion of Variance 0.9262 0.0231 0.01119 0.01107 0.00775 0.00681 0.00642
## Cumulative Proportion 0.9262 0.9493 0.96045 0.97152 0.97928 0.98609 0.99251
##              PC8      PC9      PC10
## Standard deviation  0.62065 0.60342 3.348e-15
## Proportion of Variance 0.00385 0.00364 0.000e+00
## Cumulative Proportion 0.99636 1.00000 1.000e+00
```


Do our PCA plot of this RNA-seq data

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
text(pca$x[,1], pca$x[,2], colnames(rna.data))
```



Quick barplot summary of Proportion of Variance for each PC


```
plot(pca, main="Quick scree plot")
```



Variance captured per PC

```
pca.var <- pca$sdev^2
```


Percent variance

```
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```

```
## [1] 92.6 2.3 1.1 1.1 0.8 0.7 0.6 0.4 0.4 0.0
```

Generate scree-plot

```
barplot(pca.var.per, main="Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab="Principal Component", ylab="Percent Variation")
```




Make main PCR plot more attractive and useful

```
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
     xlab=paste0("PC1 (", pca.var.per[1], "%)"),
     ylab=paste0("PC2 (", pca.var.per[2], "%)"))

text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```

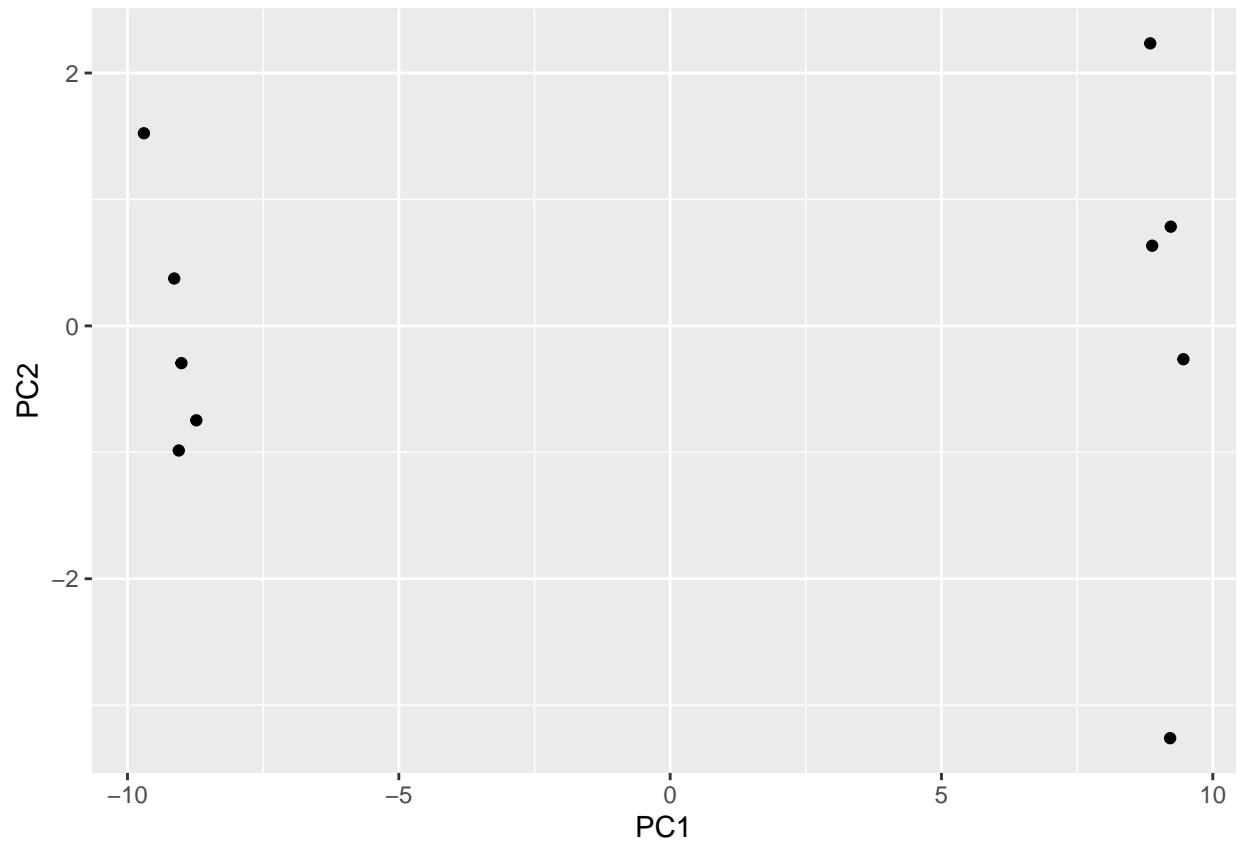


Using ggplot

```
library(ggplot2)

df <- as.data.frame(pca$x)


ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```



Add specific color and sample label aesthetics for wild-type and knock-out samples

```
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)

p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)
p
```




Add some spit and polish

```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clealy seperates wild-type from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="BIMM143 example data") +
  theme_bw()
```

PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



BIMM143 example data